

EdgeMind

Private AI that runs on your laptop, built for Snapdragon-powered HP PCs.

Ask your own files and meetings anything. Nothing leaves the device.

An on-device AI copilot that uses the Snapdragon NPU for speed, battery life and privacy.

The problem

Teams in law, healthcare and finance cannot upload sensitive documents to cloud AI. Cloud AI also adds recurring cost, needs connectivity, and running AI on the CPU drains laptop batteries.

The solution

EdgeMind indexes documents and meeting audio locally, then answers questions with cited sources, summarises long material and searches everything in one place. It works offline.

Why Snapdragon-powered HP PCs

Hexagon NPU	Embedding and speech models run on the NPU via ONNX Runtime QNN, keeping the CPU free and power low.
Power-aware policy	NPU mode and batch size follow battery state: fast when plugged in, efficient on battery.
ARM64-native and cached	No emulation overhead; compiled NPU graphs are cached for fast launches.
HP business PCs	Enterprise manageability and hardware-backed security, with MSIX/Intune deployment.

Who it is for

- Legal, healthcare and finance teams handling confidential files
- Consultants, analysts, students and researchers
- Travelling staff who need AI without a connection

Where we are

A tested prototype is ready: hardware detection, the power policy engine, an ONNX Runtime session builder with NPU, GPU and CPU fallback, a local vector store, question answering and summarisation. 7 automated tests pass. NPU benchmarks are the next step and will be measured on a Snapdragon X HP PC before any performance claim is made.

Roadmap

Weeks 1-2 prototype (done) | **3-4** NPU validation | **5-7** speech and local LLM | **8-10** desktop app | **11-12** pilot

The ask

Access to a Snapdragon X HP PC for on-device benchmarking, and a privacy-sensitive pilot team to test with.

Full details are in the accompanying proposal and presentation. Code: github.com/ayush2459/edgemind-snapdragon